

Never Let a Lead Go Cold

The first message a new lead hears back, the four guardrails that keep it from embarrassing you, and the throttle that decides how often it's allowed to speak.

A lead fills in the form on your site at 9:40 on a Tuesday night. Your workflow fires, and what it sends is "Thanks for reaching out! One of our team will be in touch within 24 to 48 hours." They read it, learn nothing, and go back to the four other tabs they had open.

Almost every follow-up build spends its effort on fixing that one message, and close to none on what the automation does afterwards. That's the wrong way round, and you can check it in your own account in about ten minutes. We read every workflow in sixteen GoHighLevel sub-accounts, 386 of them, while scoping a build for a client. All 386 were built to send. Exactly one account out of the sixteen had anything that handled a bounce, a complaint or an unsubscribe, and those three workflows had arrived inside a snapshot rather than being built by anyone who worked there. Not one account in the fleet had a rule that stopped the sending when the person wrote back.

So this kit is in two halves, and the second half is the one worth your Saturday. Part one is the message, which is a single node and a prompt you can paste. Parts two through four are the machinery around it: what happens on a reply, on a "stop", on a bounce, and on the deal that has been sitting in the same stage for a fortnight. Everything here has run in live accounts before it was written down.

PART ONE · THE FIRST MESSAGE

One **AI Agent** action, placed inside your *Website Form Leads* workflow right after whatever step first processes the form data. If you run an AI Summary step, put it after that, so the summary is available to the agent.

Connect two tools to the node, Send Email and Send SMS, and set the output mode to *No Output*. The agent writes and sends both messages itself, inside this one step. Don't add Send Email or Send SMS actions downstream, because then you have two things composing the same message and only one of them is reading the lead's own words. Both messages log to the contact's conversation history, so your team can read exactly what went out without you building a custom field to hold it.

You are {{who you are, where you are, the humans behind the business, and that you write as "we"}}.

Someone just filled out a form on your website. Send the first thing they hear back: an email first, then a text.

What they told you is in their AI Summary: [insert AI Summary field].

Their first name is {{contact.first_name}}.

Send the email with Send Email to {{contact.email}}; the text with Send SMS to {{contact.phone}}. Always send both.

THE EMAIL

Make them feel a real person read what they wrote and understands what they're after, before anyone has called. Reflect back the real thing driving them, in their own words, not a generic category. Give one useful thought tied to their actual situation that shows you know {{your field or your area}} and know what you're doing. Don't invent prices, quotes, comps, timelines or any specific you can't know. Close by letting them know {{the humans}} will be in touch personally, something to look forward to, with no timeframe and no apology for timing. Invite a reply and say it comes straight to you. Subject reads like a real reply, not marketing. Body 90 to 150 words, starts "Hi {{contact.first_name}}," and ends with {{your exact signature block}}.

THE TEXT

Two sentences, under 300 characters. Open by saying who you are, because this is a new number to them. Say you read their note and just emailed a few thoughts, and that one of you will be in touch personally. Say they can reply anytime. No subject, no signature, no timeframe.

HOW TO SOUND, both messages: warm and one to one, connected sentences with the occasional short one, plain words, their words back to them, dry warmth, no hype. Every line makes literal sense. Use contractions. Never use em dashes, "not just X but Y", buzzwords, manufactured urgency, brochure words, or exclamation-point enthusiasm.

The six slots. Who you are and where. The humans a warm lead gets handed to. Your field or service area. Your signature block. That's four you fill in once. The two that change per account are the email and phone field references, and you should pick those from the custom-value picker rather than typing them, because merge fields that look right and resolve to nothing are the most common reason a first send goes out addressed to nobody.

Build these before you turn the message on. Four small workflows, not one branching workflow with an if/else tree in it. The tree is the part that breaks without telling you, and each of these is two to four steps, so there's nothing to gain by merging them.

GUARDRAIL	TRIGGER	WHAT IT DOES
Reply, so pause and escalate	Customer replied, left unconditional on purpose	Tags the contact, tasks a human, and keeps going until somebody answers
SMS opt-out	Reply body is exactly "stop", or contains "unsubscribe", SMS only	Sets SMS DND, tags sms-opted-out
Email trouble	Mailgun reports failed, complained or unsubscribed	Sets Email DND, tags email-dnd
Stalled deal	Opportunity with no activity for 14 days, on your middle stages only	Creates a task for a human to call, and notifies

The reply guardrail is a ladder, not an alert

A notification nobody acts on is worse than no notification, because now the lead is waiting *and* your automation has stopped talking to them. So this one keeps escalating until either a human replies or the lead has at least been told a human is on the way.

WHEN	WHAT HAPPENS	ONLY IF
Straight away	Tag human-conversation-open, task for each person who could answer, notify by app and email	Always
+30 minutes	Reminder to the same people, marked urgent	No human message has gone out yet
+2 hours	The AI sends the lead a holding message, tags ai-holding-sent, notifies again	Still nothing from a human
+4 hours	Third alert, flagged as dropped	Still nothing

Those timings are a starting point rather than a rule. Set them to how fast your team can realistically answer, and decide whether the ladder holds overnight, because a 2am urgent reminder teaches people to turn notifications off.

Two details that matter more than they look. First, step one exits if `human-conversation-open` is already on the contact, so a six-message back-and-forth produces one ladder instead of six. Take the tag off when the conversation closes. Second, the holding message is allowed to acknowledge and set expectation, and nothing else. It answers only from the contact's own record and your knowledge base, and it never invents a price, a quote, a timeframe or an availability. Anything legal or financial waits for the human.

Hi Sarah, Jake and Tony have your note and one of them will come back to you personally. Wanted you to know it landed and isn't sitting in a void.

Make the alert carry what the lead said and which sequence they were in, so whoever picks it up can triage it from their phone without opening the CRM.

```
Reply from Sarah Chen (Buyer)
Sarah Chen · (503) 555-0142 · sarah@email.com
Was in: Buyer cadence, touch A3
"Yes! We'd love to see the Alberta one this weekend"
Cadence paused.
```

One more thing about this guardrail. Leave the trigger unconditional. It's tempting to filter it on message type so it only fires on SMS, and the cost of that is a channel you add in six months that silently isn't covered. A guardrail that misses a channel is worse than one that fires a little too often.

How the pause is built

The reply guardrail only tags and notifies. The stopping happens somewhere else: set `stopOnResponse` on the sending workflow itself. That's a setting on the workflow, and it's cheaper and far more reliable than having a guardrail reach in and remove contacts from other workflows by id, which is the version most people build and the version that breaks the first time a workflow gets renamed.

Why the two suppression guardrails go in first

The SMS opt-out and email trouble workflows contain no customer-facing words at all. Every step in them is either a DND action or a tag. That means there's nothing in either one to review, to get wrong, or to wait on, and their absence is a deliverability and compliance exposure the whole time they're not built. If you build one thing out of this kit tonight, build those two.

PART THREE · THE RULE THAT KEEPS THE AI OUT OF TROUBLE

An AI that has been told to be helpful and has nothing true to say will make something up, and it will do it in a confident sentence that reads beautifully. The fix isn't a better model. It's giving the automation exactly one place to get facts from, and telling it what to do when that place is empty.

Put your market claims, your seasonal facts, your current turnaround times, whatever your version of a real number is, into a **custom value** at the account level. Rewrite it on a schedule. Then write the rule into every prompt that touches it:

Facts about `{{your market or your service}}` come only from `{{custom_values.market_note}}`. If it's empty, skip this touch. Don't improvise one.

Skipping a send is a cost you can absorb. An invented figure in a message with your name at the bottom is a different kind of cost, and you find out about it later than you'd like. The same logic applies to anything you'd want a human to have chosen: if you send off-topic or seasonal messages, keep a pool of about twenty of them that your team stocks, and let the AI choose which one fits this person and write the framing. An empty pool means that message gets skipped too.

PART FOUR · THE THROTTLE

Follow-up that runs at a fixed frequency forever is how a good sequence turns into the thing people mute. Rate it against what the person is doing instead.

STATE	SIGNAL	CADENCE
Hot	Replied or clicked in the last 7 days	Three a week
Steady	Opening, not replying	Two a week
Cooling	Nothing across three touches in a row	One a week, and the next email asks directly whether this is too much
Cold	Nothing across three more	Pause, task a human, one message in 30 days

Then two standing rules on top. Every six weeks, whatever state they're in, one message offers a real way out. Every 90 days, one message rebuilds what you think they want from scratch rather than assuming the notes from March are still true. Both of those cost you a handful of unsubscribes from people who were never going to buy, and they're what stops the list rotting.

If you send SMS, cap it at one text a day inside 9am to 7pm in the contact's own timezone, and never two days running without something worth saying. Honour STOP instantly and permanently, which the opt-out guardrail already does for you.

PART FIVE · BEFORE YOU TURN IT ON

- Check that your Send Email tool lets the agent set the subject and body freely. If it's locked to a template, the prompt above can't do its job and you'll get a template with a merge field in it.
- Submit the form yourself, as a real submission, and read both messages in the contact's conversation history rather than in a preview.
- Confirm the email landed *before* the text, because the text says you just emailed them, and if the order flips you've opened with something that isn't true.
- Text yourself the word "stop" and confirm SMS DND goes on. Then check the contact record and confirm the tag is there too, because the tag is what your reporting sees.
- Read the first few real outputs out loud. The prompt holds the line, and the first real messages tell you which line it isn't holding.

Don't distribute it as a snapshot. Of the 386 workflows we read, roughly 250 were dead weight that had arrived with an account-creation snapshot and had never been opened by anyone. One six-pack of nurture workflows was sitting in draft in five different accounts and had been activated in none of them. A snapshot gives you a folder that looks like a system, and the reason nobody ever turns it on is that nobody knows what's in it. Rebuild the four guardrails in each account by hand, or compile them from a written spec. It takes an afternoon and you end up knowing what's running.

Go and read your published triggers. In that same sweep we found two live workflows firing on an unconditional "note added" trigger, meaning they ran every time anyone added a note to any contact. One of them made an AI call and wrote to the pipeline each time. The other posted the contact's details to a Zapier endpoint belonging to whoever had built the snapshot, which is data leaving the account rather than money leaving the account. Both had been running for months in an account nobody was watching. Open your workflow list, filter to published, and read what each one triggers on. It's twenty minutes and it's the highest-value twenty minutes in this document.

WHAT TO BUILD FIRST

In order: the two suppression guardrails, because they need no copy and no decisions. Then the reply ladder, once you've picked your timings. Then the first message, tested on yourself. Then the stalled-deal task. The throttle comes last, because it only matters once there's a sequence long enough to throttle.

Built by **eighty5labs**. Every pattern here ran in live GoHighLevel accounts before it was written down, and the fleet numbers come from a read-only audit of sixteen sub-accounts on 28 August 2026.

If you'd rather not build it, we do this for clients. eighty5labs.com